SABER: Support-Conditioned Adaptive Bellman Error Reweighting for Sparse HER Actor-Critic

Anonymous submission

Content Areas:
ML: Reinforcement, Imitation & Inverse RL (primary);
ROB: Robot Learning, Control & Foundation Models (secondary)

Abstract

Hindsight Experience Replay (HER) converts failed experience into goal-conditioned training data through achieved-goal relabeling. Most HER extensions alter goal selection, replay selection, or target construction. We instead ask whether per-sample squared Bellman residuals should remain equally weighted once relabeling and replay sampling are fixed. We introduce SABER (Support-Conditioned Adaptive Bellman Error Reweighting), a critic-only objective. SABER upweights a sample's squared Bellman residual when a sufficiently populated local replay set from the same discretized state–goal–horizon context contains an action in the current-policy action bin. It normalizes the weights to unit minibatch mean and uses the current-policy action only for the support query while retaining Bellman regression at the replay action. HER relabeling, uniform replay sampling, the actor objective, Bellman targets, network architectures, and the update-to-data (UTD) ratio remain unchanged. Under matched UTD-1 training on FetchPickAndPlace, SABER raises mean normalized success AUC from 0.539 to 0.609 with DDPG, from 0.372 to 0.442 with TD3, and from 0.471 to 0.583 with SAC. Paired-bootstrap 95% confidence intervals for the AUC differences exclude zero for TD3 and SAC but include zero for DDPG. On DrawerOpen with DDPG, SABER raises mean AUC from 0.207 to 0.754 and mean final success from 0.390 to 1.000. Controlled comparisons, ablations, and diagnostics examine alternative allocation mechanisms and characterize SABER's support signal. These gains, achieved under the same one-update budget and an otherwise unchanged HER pipeline, establish post-HER critic-loss allocation as a distinct design axis capable of substantial performance improvements.

1 Introduction

A central challenge in sparse-reward goal-conditioned reinforcement learning is obtaining useful supervision when successful trajectories are rare. Hindsight Experience Replay (HER) addresses this challenge by relabeling failed experience with achieved goals (Andrychowicz et al. 2017). Once

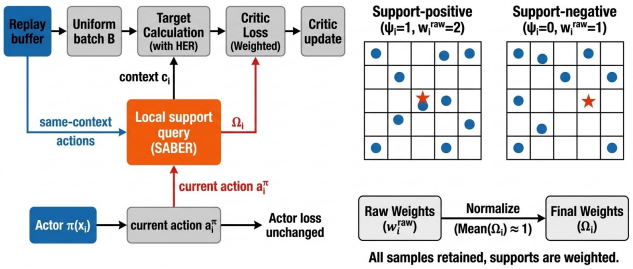


Figure 1: SABER intervenes after HER relabeling and before critic-loss aggregation. Stage 3 uses the current-policy action only to query local support; Stage 4 keeps Bellman regression at the replay action.

sampling and relabeling are fixed, however, standard HER actor–critics assign equal coefficients to all sampled Bellman residuals. This leaves a separate question: should critic supervision remain uniform when the critic regresses at replay actions but the actor queries the critic at current-policy actions?

We introduce SABER (Support-Conditioned Adaptive Bellman Error Reweighting), a critic-only post-HER objective. SABER tests whether replay actions from the same discretized state–goal–horizon context occupy the current-policy action bin, upweights support-positive residuals, and normalizes weights to unit minibatch mean. The current-policy action is used only for this support query; regression remains at the replay action. Relabeling, replay sampling, rewards, Bellman targets, actor optimization, architectures, and the update budget remain unchanged.

Our contributions are: (i) identifying post-HER critic-loss allocation as distinct from goal selection, replay prioritization, and target construction; (ii) introducing a normalized local-support objective that preserves the underlying HER actor–critic pipeline; and (iii) evaluating its cross-backbone repeatability, query and count-matched controls, update-budget trade-offs, diagnostics, and task-dependent success and failure boundaries.

2 Related Work

Goal-conditioned methods modify several stages around HER. HGG and CHER generate or schedule hindsight goals,

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.

EBP prioritizes energetic episodes, and HGR prioritizes both episodes and transition–future-goal pairs by TD error with two-stage importance correction (Ren et al. 2019; Fang et al. 2019; Zhao and Tresp 2018; Luu and Yoo 2021). MHER changes relabeled returns, ARCHER rescales hindsight rewards, and USHER corrects hindsight-induced value bias (Yang et al. 2021; Lanka and Wu 2018; Schramm et al. 2023). These methods change goal construction, experience selection, rewards, or targets. In contrast, SABER intervenes only after goal construction, replay sampling, rewards, and Bellman targets have been fixed: it preserves the relabeled minibatch and changes only the coefficient assigned to each critic Bellman residual.

General replay methods redistribute experience independently of HER. PER samples by TD error with importance correction, AER selects state-similar transitions, and ReF-ER suppresses far-policy samples while regularizing the actor (Schaul et al. 2015; **?**; Novati and Koumoutsakos 2019). SABER instead preserves minibatch membership and uses conditional action support only to set critic-loss coefficients.

Closer critic-reweighting methods learn different signals: likelihood-free importance weighting estimates a policy-to-replay density ratio, DisCor estimates target accuracy, and ReenGAGE adds goal-gradient distillation (**???**). SABER learns neither a density ratio nor an error model and adds no auxiliary target. Its distinction is a non-learned, sample-aligned current-policy action-support predicate applied to post-HER critic residuals.

3 SABER

3.1 Problem Setting and Post-HER Critic-Loss Allocation

We consider a finite-horizon goal-conditioned MDP $\mathcal{M} = (\mathcal{S}, \mathcal{A}, \mathcal{G}, P, r, \gamma, H)$, with goal-conditioned input $x = (o, g)$, policy $\pi_\theta(a \mid x)$, and critic $Q_{\phi_k}(x, a)$. Rewards are sparse: zero on goal achievement and -1 otherwise. Future-goal HER relabels a sampled transition with $\tilde{g}_i = m(s_{t'})$, $t' > t$, recomputes $\tilde{r}_i$, and produces $x_i = (o_i, \tilde{g}_i)$. DDPG, TD3, and SAC retain their standard backbone-specific targets y_i^{A}, bootstrap masks, and actor objectives (Lillicrap et al. 2015; Fujimoto, van Hoof, and Meger 2018; Haarnoja et al. 2018).

For a uniformly sampled, relabeled minibatch B, let

$$\delta_{ik} = Q_{\phi_k}(x_i, a_i^r) - y_i^{\mathsf{A}}, \qquad \mathcal{L}_Q^{\mathrm{HER}} = \frac{1}{|B|} \sum_{i \in B} \sum_{k=1}^{n_Q} \delta_{ik}^2, \tag{1}$$

where a_i^r is the replay action, $n_Q = 1$ for DDPG, and $n_Q = 2$ for TD3 and SAC. Standard HER therefore gives every sampled residual the same coefficient. SABER asks whether local replay evidence for the current-policy action a_i^π can instead allocate critic-loss weight while leaving the relabeled batch and y_i^{A} fixed.

3.2 Current-Policy Local Support

For sample i, SABER builds a task-specific discrete context $c_i = C(o_i, \tilde{g}_i, h_i)$, where h_i is the remaining horizon. A bounded replay index retrieves at most N_{scan} same-coarse-context candidates and filters them by the full context, yielding the local index set I_i. The relabeled query goal $\tilde{g}_i$ is used for both query and candidate contexts. Exact context features and indexing details are provided in the supplementary artifact.

Let $b_K(a)$ denote the multidimensional bin obtained by normalizing each action dimension to its environment bounds and dividing it into K bins, and let $\mathcal{U}_i = \{b_K(a_j^r) : j \in I_i\}$. With $n_i = |I_i|$, SABER defines

$$\psi_i = \mathbf{1}\{n_i \geq N_{\mathrm{local}}\}\, \mathbf{1}\{\min(n_i, M_{\mathrm{cap}}) \geq M_{\mathrm{avail}}\} \cdot \mathbf{1}\{b_K(a_i^\pi) \in \mathcal{U}_i\}\,. \tag{2}$$

The main setting uses $K = 5$, $N_{\mathrm{scan}} = 512$, $N_{\mathrm{local}} = 4$, $M_{\mathrm{cap}} = 16$, and $M_{\mathrm{avail}} = 4$. DDPG and TD3 use their deterministic actor action as a_i^π, while SAC uses its deterministic squashed mean. The indicator uses neither Q-values nor TD errors: the current-policy action is used only for the discrete support query, and Bellman regression remains at a_i^r.

3.3 Normalized Critic-Only Objective

SABER assigns and normalizes per-sample weights as

$$w_i^{\mathrm{raw}} = 1 + \lambda_S \psi_i, \qquad \bar{w}_B = \frac{1}{|B|} \sum_{j \in B} w_j^{\mathrm{raw}}, \tag{3}$$

$$\Omega_i = \frac{w_i^{\mathrm{raw}}}{\max(\epsilon_w, \bar{w}_B)}. \tag{4}$$

The main method uses $\lambda_S = 1$, giving raw weights two and one to support-positive and support-negative samples. Its critic objective is

$$\mathcal{L}_Q^{\mathrm{SABER}} = \frac{1}{|B|} \sum_{i \in B} \Omega_i \sum_{k=1}^{n_Q} \delta_{ik}^2. \tag{5}$$

For TD3 and SAC, the same Ω_i weights both critic residuals. All coefficients remain positive and have minibatch mean one, so SABER preserves minibatch membership and approximately preserves the per-update loss scale while changing the relative residual contributions.

The main variant leaves the backbone actor objective unchanged. This preserves the DDPG, TD3, and SAC policy-update rules and isolates Bellman-loss allocation; critic-only is attribution-focused rather than claimed to be performance-optimal. Overall, SABER leaves HER, replay sampling, Bellman targets, architectures, optimizers, exploration, and UTD unchanged. Support is a binary occupancy proxy, not an estimate of density, uncertainty, or target accuracy, and SABER does not directly regress at the current-policy action or generate actions absent from replay.

4 Experiments

4.1 Common Experimental Protocol

Tasks, cohorts, and training protocol. We evaluate FetchPickAndPlace-v4 (FPP; horizon 50, 1M environment steps) with DDPG, TD3, and SAC, and Meta-World

DrawerOpen and DrawerClose (horizon 200, 2M steps) with DDPG and TD3. Primary comparisons use paired seeds 1–5, future-goal HER probability 0.8, and UTD 1. We make no paired claims for SAC on the Drawer tasks because no complete matched HER–SABER cohort exists: the available DrawerOpen SAC+HER runs remain at zero success for all five seeds, and neither matched SABER runs nor DrawerClose SAC results were collected. Sparse rewards and future relabeling follow HER (Andrychowicz et al. 2017); backbone targets and policy updates follow DDPG, TD3, and SAC (Lillicrap et al. 2015; Fujimoto, van Hoof, and Meger 2018; Haarnoja et al. 2018). Shared continuous-control settings are fixed within each matched task–backbone cohort without method-specific retuning; complete values, implementation details, and provenance are provided in the supplement.

Evaluation and statistics. Policies are evaluated deterministically over 20 episodes at each checkpoint. Paired analyses use only exact shared evaluation steps, without smoothing, interpolation, or filling. Normalized success AUC is the trapezoidal area divided by the shared step span; final success and Hit90 (the first checkpoint with success ≥ 0.9, reported with reached-seed coverage) are secondary. We report seed means and sample standard deviations; paired claims use within-seed differences and 10,000-resample percentile-bootstrap 95% confidence intervals. Absolute AUC values are not compared across tasks with different horizons or budgets.

4.2 Primary Comparisons

The primary comparison matches future-goal HER against HER+SABER, changing only critic-loss coefficients. We then compare intervention families. Full HGR prioritizes episodes and transition–future-goal pairs by TD error with two-stage importance correction (Luu and Yoo 2021). PER-style retains uniform replay but uses normalized, capped TD-error critic weights, whereas Full PER changes replay-slot sampling and applies annealed importance correction (Schaul et al. 2015). Where reported, these cohorts match task, backbone, seeds, training budget, and UTD; method-specific priority and correction settings are given in the supplement. The no-HER ER/AnnealingER/AER block, high-replay-ratio HER, and REDQ-HER change the learning pipeline, update budget, or architecture and are therefore reported only as secondary scaling references.

4.3 Mechanism and Weight-Scope Ablations

Using DDPG on FPP, DrawerOpen, and DrawerClose, we isolate support-query identity, sample-wise weight assignment, and optimization scope. Variants use paired seeds 1–5 and preserve the task-specific budget, HER and replay procedures, Bellman targets, architecture, optimizer settings and base learning rates, exploration, UTD, and evaluation grid. Within a task, support hyperparameters are not retuned.

Query and assignment controls. Replay-Action Query substitutes a_i^r for a_i^π; Shuffled-Policy Query applies a within-minibatch derangement, preserving the action multiset while disrupting sample–query alignment. These perturbations can also change the positive count, so assignment controls handle that confound separately. At each such control update, Main's indicator rule is applied to that control run's own minibatch and replay state, giving $k = \sum_i \psi_i^{\mathrm{Main}}$. Random-Support samples k batch positions without replacement, matching the count and weight histogram but removing sample-level support assignment. Strict Context-Eligible Count-Matched instead samples k positions without replacement from the same eligibility set E; its only intended remaining difference from Main is membership-aligned location.

Optimization scope. Critic-Only applies Ω_i to Bellman losses, Actor-Only applies it only to the actor objective, and Actor+Critic applies it to both. The query rule, support hyperparameters, and normalization rule are fixed across scopes, and Ω_i is always stop-gradient. Critic-Only is the attribution-focused Main, not an assumed performance optimum.

Diagnostic instrumentation. SABER runs additionally record support activation, absolute TD error, and same-minibatch residual changes without adding gradients or altering replay, objectives, or update order. FPP diagnostics use five seeds per backbone. DrawerOpen DDPG uses a separate ten-seed diagnostic cohort: all ten enter its primary diagnostic aggregate, while the original matched five remain the primary performance cohort and outcome-stratified summaries are secondary. Measurements are averaged within seed before aggregation; field and phase definitions appear in the supplement.

5 Results

5.1 Cross-Backbone Repeatability

SABER produces a repeatable cross-backbone improvement under matched one-update budgets.

SABER raises mean FPP AUC from 0.539 to 0.609 on DDPG, 0.372 to 0.442 on TD3, and 0.471 to 0.583 on SAC. The paired AUC intervals exclude zero for TD3 and SAC, and all three backbones improve in mean final success and Hit90 coverage. This demonstrates cross-backbone compatibility rather than a single-backbone anomaly.

Ten-seed TD3 confirmation. A frozen paired extension on seeds 6–10 gives ten matched TD3 pairs on the exact 10k–1M grid. HER versus SABER reaches AUC 0.383 ± 0.041 versus 0.429 ± 0.058, with paired delta $+0.0467 [0.0085, 0.0783]$; 8/10 seed deltas are positive and every leave-one-out mean remains positive. The extension corroborates the five-seed result while preserving its finite-seed boundary: DDPG supplies mean-improvement rather than interval-closed evidence, and neither result implies improvement on every seed.

5.2 Policy-Aligned Querying Is Substantive

The support query is not interchangeable with replay-action querying or sample-misaligned policy querying.

The main current-policy query exceeds replay-query and shuffled-policy on every paired seed. Replay-query has a higher support-positive rate than Main (0.426 versus 0.219) yet lower AUC, while shuffling preserves the policy-action marginal distribution but removes sample alignment. Thus query identity and sample alignment matter beyond merely producing more support-positive weights.

Table 1: Ablations targeting query identity, weight assignment, and optimization scope.

Axis	Variant	Role or controlled question
Query	Replay-Action Query	Effect of replacing a_i^π with a_i^r
Query	Shuffled-Policy Query	Effect of disrupting sample–query alignment
Assignment	Random-Support (Batch-Wide Count Matched)	Effect of preserving the per-update positive count and weight histogram while randomizing sample-level assignment
Assignment	Strict Context-Eligible Count-Matched	Membership effect beyond positive count and context eligibility
Scope	Critic-Only (Main)	Reference intervention applying Ω_i only to Bellman losses
Scope	Actor-Only	Comparison of actor-only with critic-only weighting
Scope	Actor+Critic	Effect of adding actor weighting to critic weighting

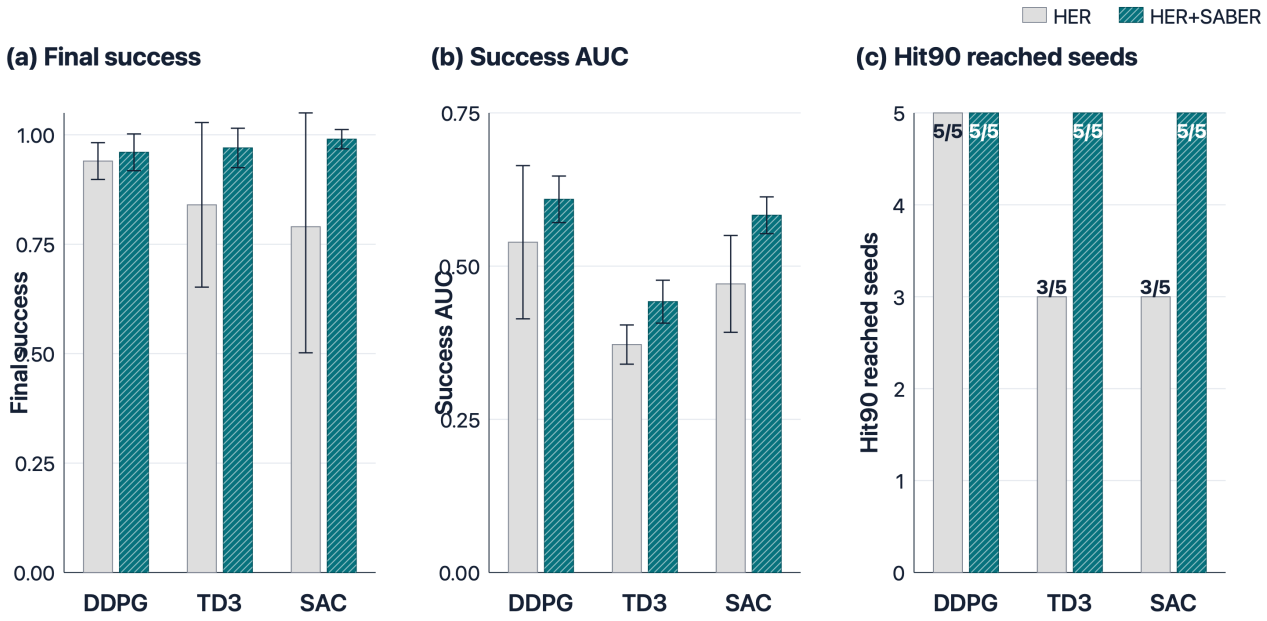


Figure 2: FetchPickAndPlace evidence across matched one-update-per-step HER backbones. Bars and error bars show the seed mean and standard deviation. Exact paired 95% SABER–HER intervals are reported in Table 2; the AUC interval includes zero for DDPG and excludes zero for TD3/SAC.

Count-matched assignment. Each control update recomputes E, M, k and samples $Z \subseteq E$, $|Z| = k$, without using membership locations; instrumented summaries record no count or weight-distribution violations. Main versus control AUC is 0.4424 ± 0.0350 versus 0.3908 ± 0.0700, with paired difference $+0.0516[-0.0165, 0.0992]$, 4/5 positive seeds, and exact $p = 0.25$. This removes the count confound and directionally favors membership alignment, but does not confirm its independent contribution.

5.3 A Favorable Gradient-Update Trade-Off

On DDPG and TD3, SABER reaches AUC comparable to $4\times$ replay-ratio HER with one quarter of the gradient-update budget.

For TD3, comparable AUC is accompanied by substantially higher final success (0.970 versus 0.610) and Hit90 coverage (5/5 versus 1/5). SABER therefore offers a complementary one-update allocation strategy rather than relying on brute-force replay scaling. **Boundary:** this is a gradient-update trade-off, not a wall-clock-efficiency claim; SAC high-RR and REDQ-HER TQB attain higher AUC with larger update or architecture budgets.

5.4 Large-Effect Rescue on DrawerOpen

SABER rescues a failure-prone HER regime on DrawerOpen. It raises final success from 0.390 to 1.000, AUC from 0.207 to 0.754, and Hit90 coverage from 2/5 to 5/5 seeds. This shows that critic-side allocation can have a

Table 2: Matched-backbone comparison: FPP main results and matched HGR reference at 1M steps. SAC AUC and Hit90 use the exact shared 50k–1M grid. SABER improves means across all backbones; paired AUC CIs are clearer for TD3/SAC than DDPG.

Backbone	Method	Final	AUC	Hit90	Steps/sec
DDPG	HER	0.940 ± 0.042	0.539 ± 0.125	600k (5/5)	65.03
DDPG	**HER+SABER (Ours)**	$\mathbf{0.960 \pm 0.042}$	$\mathbf{0.609 \pm 0.038}$	**540k (5/5)**	**41.72**
DDPG	HGR (matched)	0.540 ± 0.365	0.338 ± 0.149	930k (2/5)	16.93
TD3	HER	0.840 ± 0.188	0.372 ± 0.032	980k (3/5)	65.33
TD3	**HER+SABER (Ours)**	$\mathbf{0.970 \pm 0.045}$	$\mathbf{0.442 \pm 0.035}$	**840k (5/5)**	**41.25**
SAC	HER	0.790 ± 0.288	0.471 ± 0.079	800k (3/5)	38.93
SAC	**HER+SABER (Ours)**	$\mathbf{0.990 \pm 0.022}$	$\mathbf{0.583 \pm 0.030}$	**750k (5/5)**	**29.10**
DDPG	Δ SABER–HER	$+0.020[-0.040, 0.070]$	$+0.070[-0.016, 0.193]$	–	–
TD3	Δ SABER–HER	$+0.130[-0.030, 0.310]$	$+0.071[0.032, 0.095]$	–	–
SAC	Δ SABER–HER	$+0.200[-0.020, 0.430]$	$+0.112[0.064, 0.161]$	–	–

Table 3: FPP TD3 mechanism controls at 1M steps, seeds 1–5. Paired 95% CIs compare Main with each row. The strict control samples $k = |M|$ positions from E on each control update; context-count only is count-unmatched.

Query / control	Final	AUC	Hit90	Main − row AUC [CI]
HER	0.840 ± 0.188	0.372 ± 0.032	980k (3/5)	$+0.071[0.032, 0.095]$
Main SABER (Ours)	$\mathbf{0.970 \pm 0.045}$	$\mathbf{0.442 \pm 0.035}$	**840k (5/5)**	**–**
Replay action	0.660 ± 0.219	0.368 ± 0.042	920k (1/5)	$+0.075[0.059, 0.092]$
Shuffled policy	0.550 ± 0.226	0.347 ± 0.051	890k (1/5)	$+0.096[0.063, 0.128]$
Context-eligible count-matched	0.910 ± 0.096	0.391 ± 0.070	955k (4/5)	$+0.052[-0.016, 0.099]$
Context-count only	0.800 ± 0.242	0.396 ± 0.090	850k (3/5)	$+0.047[-0.041, 0.122]$

Table 4: Scaling references: FPP SAC/REDQ metrics use the exact shared 50k–1M grid. High-RR and REDQ rows expose update/backbone budget tradeoffs.

Backbone	Method	Updates/step	Final	AUC	Hit90	Runtime h
DDPG	**HER+SABER (Ours)**	**1**	$\mathbf{0.960 \pm 0.042}$	$\mathbf{0.609 \pm 0.038}$	**540k (5/5)**	**6.66**
DDPG	HER high-RR $4\times$	4	0.990 ± 0.022	0.618 ± 0.071	480k (5/5)	16.11
DDPG	HER high-RR $8\times$	8	0.720 ± 0.407	0.443 ± 0.269	620k (3/5)	31.85
TD3	**HER+SABER (Ours)**	**1**	$\mathbf{0.970 \pm 0.045}$	$\mathbf{0.442 \pm 0.035}$	**840k (5/5)**	**6.74**
TD3	HER high-RR $4\times$	4	0.610 ± 0.243	0.449 ± 0.079	630k (1/5)	8.76
SAC	**HER+SABER (Ours)**	**1**	$\mathbf{0.990 \pm 0.022}$	$\mathbf{0.583 \pm 0.030}$	**750k (5/5)**	**9.55**
SAC	HER high-RR $4\times$	4	0.960 ± 0.042	0.768 ± 0.097	450k (5/5)	47.02
REDQ-SAC	REDQ-HER TQB	4	1.000 ± 0.000	0.867 ± 0.052	300k (5/5)	89.52

large practical effect when matched HER suffers from seed-level failures.

Random-support and PER-style controls are also strong, so DrawerOpen establishes intervention-family significance rather than unique structured-support causality. The HGR-style sampler-only row reaches Hit90 in 2/5 seeds; SABER minus HGR-style is $+0.443[0.106, 0.780]$ in AUC, but the row is not full HGR. Scaled REDQ-HER TQB reaches Hit90 earlier (370k versus 500k median), whereas SABER has higher mean AUC and lower variance, 0.754 ± 0.044 versus 0.672 ± 0.314; neither method dominates every seed or metric.

An independent SABER-only extension reaches Hit90 in 4/5 additional seeds; descriptively pooled over seeds 1–10, SABER reaches 9/10 with final success 0.895 ± 0.315 and AUC 0.652 ± 0.246. Because no matched methods were extended, the original five-seed paired comparisons remain primary. A separate diagnostic rerun reproduces the original five-seed aggregate AUC (0.755 versus 0.754). Across nine successful diagnostic seeds, support-positive samples show $1.17\times$ pre-update TD error and $2.1\times$ same-minibatch residual reduction; failed seed 7 first loses contact exploration and later exhibits support saturation. These in-sample diagnostics characterize allocation and failure behavior rather than causal generalization.

Table 5: DrawerOpen stress case at 2M steps. DDPG rows use the local matched regime; REDQ-HER TQB is an N=10, M=2, UTD-4 architecture/compute-scaled reference. Hit90 reports median first hit and reached seeds.

Method	Final	AUC	Hit90
HER	0.390 ± 0.534	0.207 ± 0.287	960k (2/5)
HGR-style sampler-only	0.400 ± 0.548	0.311 ± 0.425	450k (2/5)
HER+SABER (Ours)	$\mathbf{1.000 \pm 0.000}$	$\mathbf{0.754 \pm 0.044}$	**500k (5/5)**
Random-support	1.000 ± 0.000	0.706 ± 0.122	510k (5/5)
PER-style	0.800 ± 0.447	0.715 ± 0.090	610k (5/5)
REDQ-HER TQB (scaled)	1.000 ± 0.000	0.672 ± 0.314	370k (5/5)

Table 6: Control closure: what the reported controls establish and what remains unresolved.

Evidence	Established	Remaining interpretation
DDPG controls: Random-fixed AUC 0.540 is close to HER 0.539; Random-support AUC 0.604 is close to SABER 0.609.	Fixed-count random weighting alone does not explain the gain.	FPP DDPG does not isolate exact sample-level alignment; count-preserving randomization can regularize.
TD3/SAC random controls stay near or below SABER in AUC.	Random reweighting does not explain all gains across backbones.	TD3/SAC provide the cleanest random-control separation.
High-RR/REDQ rows match or exceed SABER under larger budgets.	SABER offers a distinct one-update allocation strategy.	It does not dominate update or architecture scaling.
FPP HGR (matched) uses pair priorities and episode/pair IS correction but reaches AUC 0.338 and Hit90 in 2/5; SABER reaches 0.609 and 5/5.	Matched goal prioritization does not explain SABER's gain in this regime.	This is not an exact original-paper reproduction or a general HGR ranking.
Diagnostics show weak TD-error correlation, top-error enrichment, and larger immediate correction in TD3/SAC.	SABER is not PER-style largest-error selection.	The evidence supports error allocation, not lower target variance.
Main query AUC 0.442 exceeds replay 0.368 and shuffled 0.347 on all five paired seeds; the context-eligible count-matched control reaches 0.391.	Current-policy identity and sample alignment are substantive, and the strict control removes the positive-count/context-eligibility confound.	Main–control is $+0.052[-0.016, 0.099]$ on 4/5 seeds; membership-specific benefit is directional but not interval-closed.

5.5 Mechanism Characterization and Boundaries

Allocation diagnostics. Instrumented DDPG, TD3, and SAC runs log 500 full diagnostic batches and 64k sample rows per backbone. Support-positive rates are 0.367, 0.219, and 0.197, while Pearson correlations with $|\delta|$ are only 0.116, 0.128, and 0.143. Their top-10% TD-error overlaps are 0.516, 0.360, and 0.342: enriched over each support base rate but far from TD-error equivalence. TD3/SAC support-positive samples also receive larger immediate residual correction after weight normalization. Target variance is higher, not lower, in all three diagnostics, supporting an allocation signal rather than a lower-noise interpretation.

Together, the controls establish a distinct policy-aligned allocation signal on TD3/SAC and characterize where generic reweighting, goal prioritization, or scaling remain competitive. They define the operating regime of the contribution rather than displacing its core matched-budget findings.

6 Discussion

6.1 What SABER Establishes

SABER establishes post-HER critic supervision allocation as a distinct design axis. Holding relabeling, replay, actor optimization, Bellman targets, critic architecture, and update rate fixed provides cleaner effect attribution and a common interface across DDPG, TD3, and SAC. The evidence then localizes the signal: aligned current-policy querying exceeds replay and shuffled-policy queries, cross-backbone and ten-seed results show repeatability, DDPG/TD3 expose a favorable gradient-update trade-off, and DrawerOpen demonstrates that critic-side allocation can rescue a failure-prone HER regime.

6.2 Scope and Limitations

The evidence is scoped in five ways. First, the support predicate is task-specific and uses coarse membership; FetchSlide DDPG decreases from HER final/AUC $0.630/0.324$ to SABER $0.590/0.301$, and DrawerClose remains high variance. Second, query and membership assignment are tested only on FPP TD3. The strict context-eligible count-matched control removes the earlier count confound and favors Main in mean AUC, but its five-seed interval includes zero; a weakly active D8 gate also does not establish a diversity benefit. Third, diagnostics are in-sample allocation measurements, not lower-noise or causal-generalization proofs.

Fourth, the current lookup is costly: in a hardware-specific seed-1 profile, TD3 throughput falls from 124.0 to 41.4 steps/s and support/weighting occupies 62.3% of SABER's loop. Fifth, matched FPP HGR is not an exact original-paper reproduction, HGG/MHER/CHER remain untested, and most controls use five seeds.

Statistical scope. Paired deltas use 10,000 seed bootstrap resamples. Exact sign-flip tests for FPP DDPG/TD3/SAC AUC give $p = 0.375/0.125/0.0625$, with positive deltas on 3/5, 4/5, and 5/5 seeds; all leave-one-out means remain positive. With five nonzero pairs, 0.0625 is the minimum attainable two-sided value. The ten-seed TD3 extension contains two negative deltas and exact $p = 0.0527$, while the DrawerOpen extension is SABER-only and unmatched. These checks establish repeatability and outlier stability, not improvement on every seed.

7 Conclusion

SABER establishes post-HER critic supervision allocation as a design axis distinct from relabeling, replay, actor optimization, Bellman targets, architecture, and update rate. It yields a repeatable cross-backbone FPP signal, corroborated by the paired TD3 ten-seed extension, and aligned current-policy querying outperforms replay/shuffled alternatives. Count matching removes the earlier confound but does not interval-close the incremental membership effect. SABER also offers a DDPG/TD3 gradient-update trade-off and rescues a failure-prone DrawerOpen regime. Generic support representations remain future work.

8 Reproducibility and Anonymity

An anonymized archive provides sanitized configs, exact-grid curves, package-local analyzers/tests, and SHA256 manifests for the matched HGR, TD3 ten-seed, query/count-matched, DrawerOpen, scaling, and runtime evidence. Package commands regenerate compact tables; independent checks cover exact inference, leave-one-out, and table consistency. Public URLs, user namespaces, and raw commit identities are excluded.

References

Andrychowicz, M.; Wolski, F.; Ray, A.; Schneider, J.; Fong, R.; Welinder, P.; McGrew, B.; Tobin, J.; Abbeel, P.; and Zaremba, W. 2017. Hindsight Experience Replay. *Advances in Neural Information Processing Systems*.

Fang, M.; Zhou, T.; Du, Y.; Han, L.; and Zhang, Z. 2019. Curriculum-Guided Hindsight Experience Replay. In *Advances in Neural Information Processing Systems*, 12623–12634.

Fujimoto, S.; van Hoof, H.; and Meger, D. 2018. Addressing Function Approximation Error in Actor-Critic Methods. *Proceedings of the 35th International Conference on Machine Learning*.

Haarnoja, T.; Zhou, A.; Abbeel, P.; and Levine, S. 2018. Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor. *Proceedings of the 35th International Conference on Machine Learning*.

Lanka, S.; and Wu, T. 2018. ARCHER: Aggressive Rewards to Counter Bias in Hindsight Experience Replay. *arXiv preprint arXiv:1809.02070*.

Lillicrap, T. P.; Hunt, J. J.; Pritzel, A.; Heess, N.; Erez, T.; Tassa, Y.; Silver, D.; and Wierstra, D. 2015. Continuous Control with Deep Reinforcement Learning. *arXiv preprint arXiv:1509.02971*.

Luu, T. M.; and Yoo, C. D. 2021. Hindsight Goal Ranking on Replay Buffer for Sparse Reward Environment. *IEEE Access*.

Novati, G.; and Koumoutsakos, P. 2019. Remember and Forget for Experience Replay. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, 4851–4860.

Ren, Z.; Dong, K.; Zhou, Y.; Liu, Q.; and Peng, J. 2019. Exploration via Hindsight Goal Generation. In *Advances in Neural Information Processing Systems*.

Schaul, T.; Quan, J.; Antonoglou, I.; and Silver, D. 2015. Prioritized Experience Replay. *arXiv preprint arXiv:1511.05952*.

Schramm, L.; Deng, Y.; Granados, E.; and Boularias, A. 2023. USHER: Unbiased Sampling for Hindsight Experience Replay. In *Proceedings of the 6th Conference on Robot Learning*, volume 205 of *Proceedings of Machine Learning Research*, 2073–2082.

Yang, R.; Lyu, J.; Yang, Y.; Yan, J.; Luo, F.; Luo, D.; Li, L.; and Li, X. 2021. Bias-Reduced Multi-Step Hindsight Experience Replay for Efficient Multi-Goal Reinforcement Learning. *arXiv preprint arXiv:2102.12962*.

Zhao, R.; and Tresp, V. 2018. Energy-Based Hindsight Experience Prioritization. In *Proceedings of the Conference on Robot Learning*, 113–122.